

Return oriented programming (ROP)

Cat and mouse

- **Defense:** Make stack/heap nonexecutable to prevent injection of code
 - **Attack response:** Jump/return to libc
- **Defense:** Hide the address of desired libc code or return address using ASLR
 - **Attack response:** Brute force search (for 32-bit systems) or **information leak** (format string vulnerability)
- **Defense:** Avoid using libc code entirely and use code in the program text instead
 - **Attack response:** Construct needed functionality using **return oriented programming (ROP)**

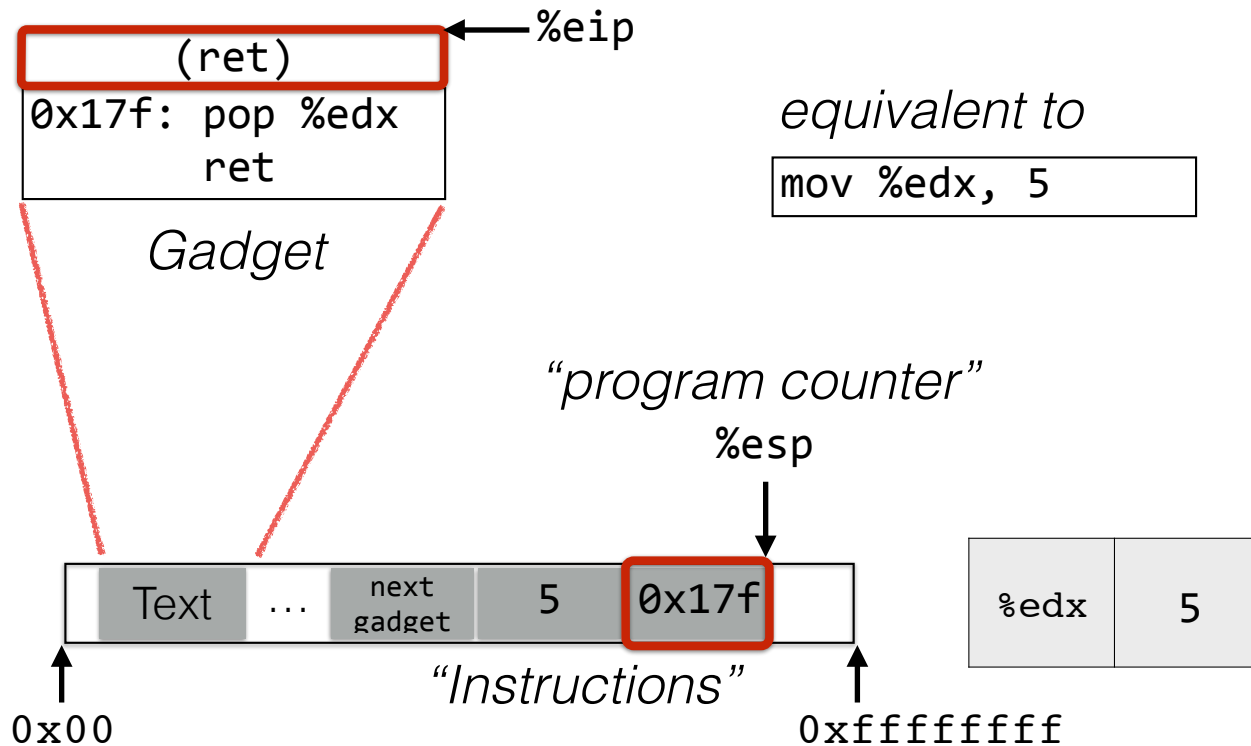
Return-oriented Programming

- Introduced by Hovav Shacham in 2007
 - *The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)*, CCS'07
- Idea: rather than use a single (libc) function to run your shellcode, **string together pieces of existing code, called *gadgets***, to do it instead
- Challenges
 - **Find the gadgets** you need
 - **String them together**

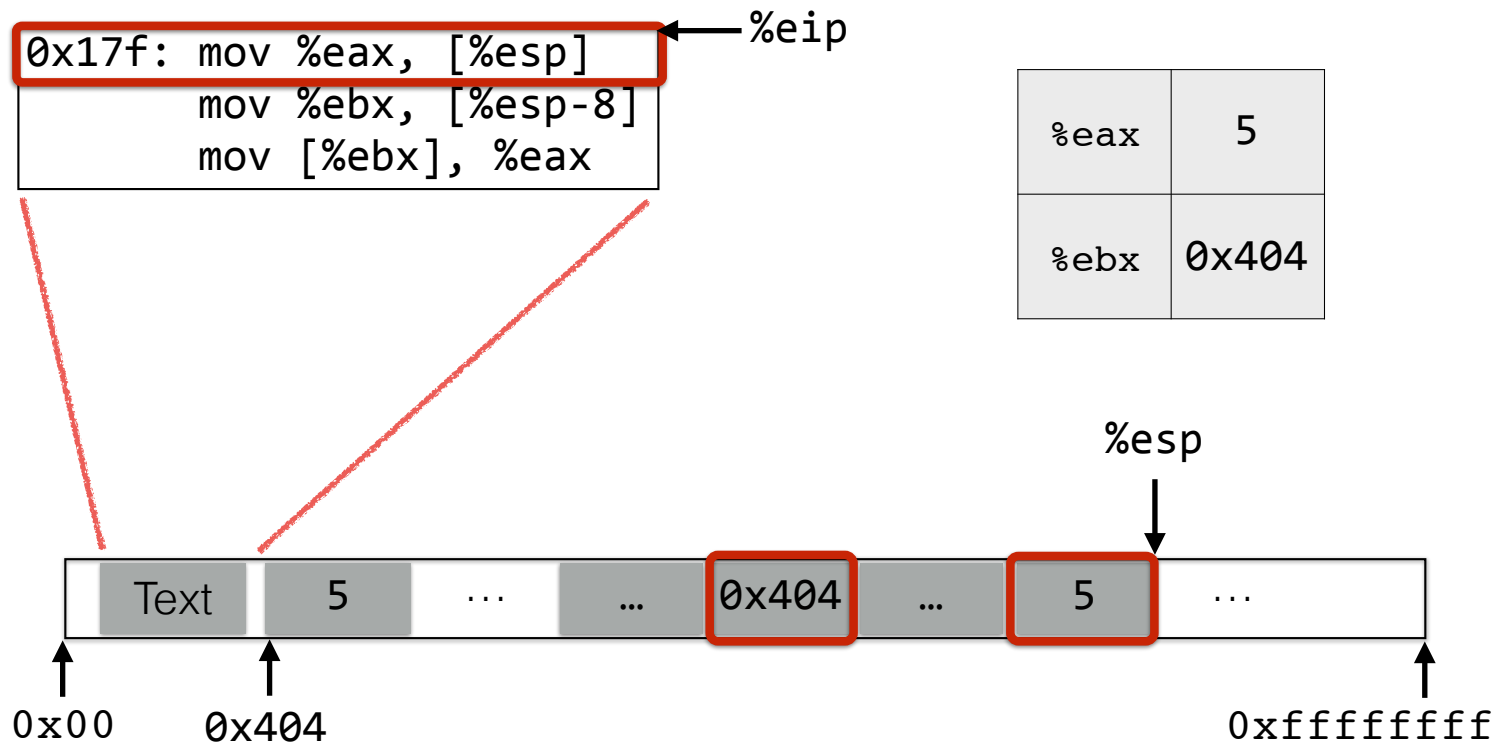
Approach

- Gadgets are instruction groups that end with `ret`
- Stack serves as the code
 - `%esp` = program counter
 - Gadgets invoked via `ret` instruction
 - Gadgets get their arguments via `pop`, etc.
 - Also on the stack

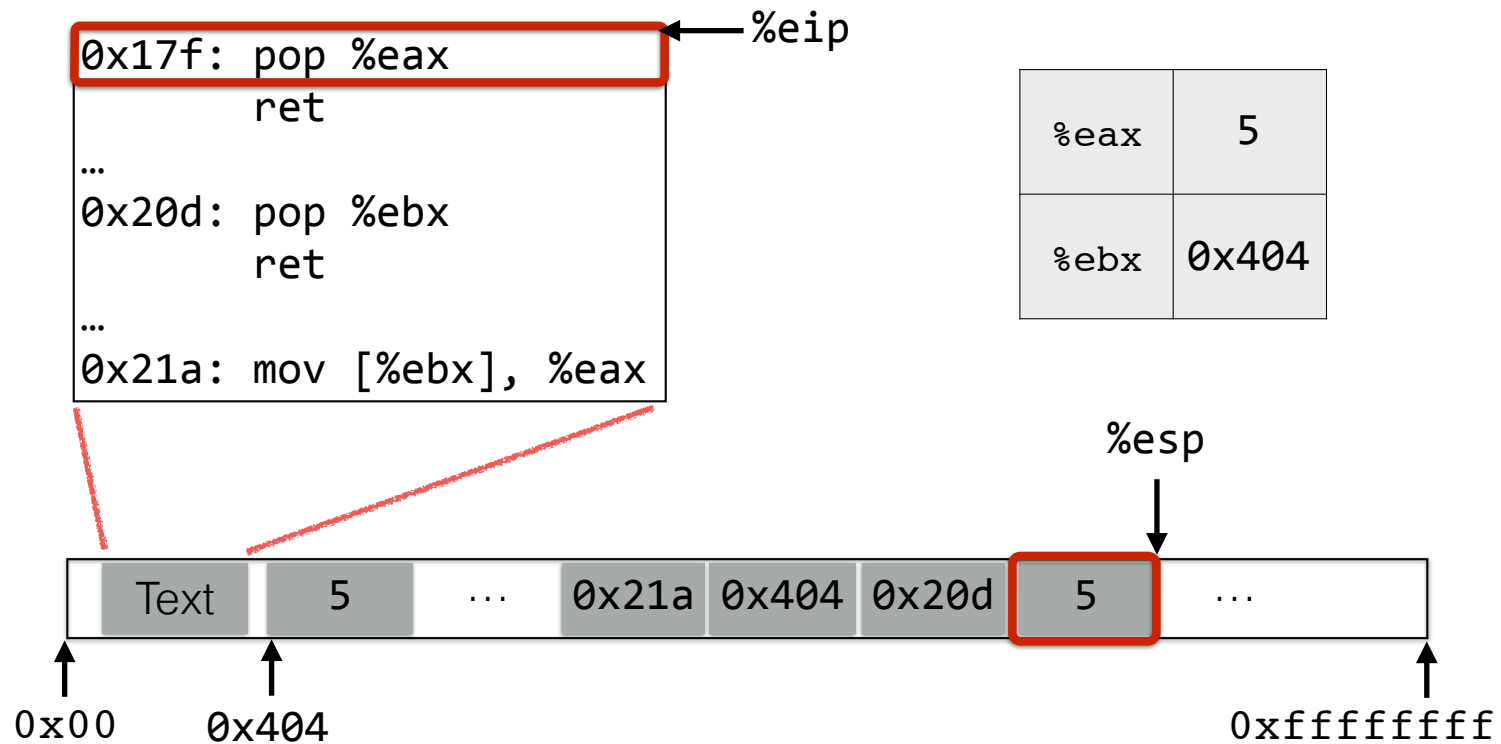
Simple example



Code sequence



Equivalent ROP sequence



Return-Oriented Programming

is A lot like a ransom
note, BUT instead of cutting
out Letters from Magazines,
YOU ARE cutting out
instructions from text
segments

Image by Dino Dai Zovi

Whence the gadgets?

- How can we find gadgets to construct an exploit?
 - **Automate a search of the target binary for gadgets**
(look for `ret` instructions, work backwards)
 - Cf. <https://github.com/0vercl0k/rp>
- Are there sufficient gadgets to do anything interesting?
 - Yes: Shacham found that for significant codebases (e.g., `libc`), **gadgets are Turing complete**
 - Especially true on x86's dense instruction set
 - Schwartz et al (USENIX Security '11) have automated gadget shellcode creation, though not needing/ requiring Turing completeness

Blind ROP

- **Defense: Randomizing the location of the code**
(by compiling for position independence) on a 64-bit machine makes attacks very difficult
 - Recent, published attacks are often for 32-bit versions of executables
- **Attack response: Blind ROP**
 - If server restarts on a crash, but does not re-randomize:
 - 1. Read the stack to **leak canaries and a return address**
 - 2. Find gadgets (at run-time) to **effect call to write**
 - 3. **Dump binary to find gadgets for shellcode**

<http://www.scs.stanford.edu/brop/>

Defeat!

- The blind ROP team was able to **completely automatically**, only **through remote interactions**, develop a **remote code exploit for nginx**, a popular web server
 - The exploit was carried out on a 64-bit executable with full stack canaries and randomization
- Conclusion: **give an inch, and they take a mile?**
- Put another way: **Memory safety is really useful!**