

This document describes what the vsftpd code trusts, what it doesn't trust, and the reasoning behind any trust decisions.

#### The importance of trust and trust relationships

=====

Imagine a largely well written and secure piece of code. Now imagine that this piece of code delegates a task to an external program, perhaps in the name of code reuse. Now, if this external program is sloppily coded and insecure, we've wasted a lot of effort making our original program secure; our erroneous trust of the buggy external program means we have a security leak, even though we were careful in `_our_` code.

There is a very similar situation with buggy library APIs. Imagine our secure program calling some complex library function which lets the side down by containing a security hole.

Lets put some concrete examples on the two similar above considerations. We can even give examples in the context of FTP daemons.

##### 1) External `/bin/ls` helper

A very common operation asked of FTP servers is to provide a directory listing. Unfortunately, convention seems to be to emit the directory listing in UNIX `"/bin/ls -l"` format. Even the Microsoft FTP service can be observed to do this. When writing an FTP server for the UNIX platform, then, this leads to the temptation to reuse `/bin/ls` as a child process, to avoid having to rewrite a load of code to handle directory listings.

Even more unfortunately, FTP server writers seem to want to adopt the versatility of the average `/bin/ls` implementation. This means they allow clients to specify arbitrary parameters to `/bin/ls`.

By using an external `/bin/ls` command, we would tie the security of our FTP server to that of the `/bin/ls` code. Be careful not to underestimate the amount of code paths in `/bin/ls` which are explorable by a remote malicious user. GNU `/bin/ls` has a myriad of options. Some of these options are complex such as `-I` or the various formatting options. All it takes is a single coding flaw in the handling of one of these options, and your FTP security is in trouble.

By using an external `/bin/ls`, you also inherit the risk of any dangerous or complex APIs it uses. For example, calls to libc's complex `fnmatch()` or `glob()` functions, which will get given arbitrary malicious user controlled

data as the search patterns. Also remember that users (and sometimes remote users) can upload/create files, and filenames are a very prominent input to /bin/ls.

To conclude: vsftpd has no intention of using an external /bin/ls program because of the risks outlined above. Even if I were to audit e.g. GNU fileutils /bin/ls, and also important parts of glibc, this would still leave security in an unknown state on other platforms. The solution I have employed is to write a minimal internal implementation of a /bin/ls listing generator; it's hardly difficult. As a happy side effect, this will boost performance by avoiding unnecessary fork()s and exec()s!

Here's some quick data about FTP servers which tend to use external ls programs:

```
ftp.wuftp.org:
ftp> ls --version
227 Entering Passive Mode (x.x.x.x.x)
150 Opening ASCII mode data connection for /bin/ls.
ls (GNU fileutils) 3.16
226 Transfer complete.
```

```
ftp.digital.com:
ftp> ls -v
227 Entering Passive Mode (x.x.x.x.x)
150 Opening ASCII mode data connection for /bin/ls.
/bin/ls: illegal option -- v
usage: ls [ -lACFLRabcdfgilmnopqrstux ] [files]
226 Transfer complete.
```

Note that /bin/ls is not the only external program invoked by common FTP servers such as wu-ftp. wu-ftp also has the ability to invoke "tar" and "gzip" on the fly, so there are trust relationships there too.

## 2) Complex library APIs

vsftpd is very careful to avoid using library calls which are potentially dangerous. I would typically classify calls as dangerous if they interact with the network non-trivially, or take malicious user supplied data and start parsing it in a major way.

Some examples are clearly required (vsftpd avoids using any of the following):

1) `fnmatch()`. This is the `libc` glob pattern matcher. The danger comes from the fact that the user supplies the glob pattern - "`ls *.mp3`" would be a simple example. Furthermore, glob pattern matching is complex and involves a lot of string handling.

2) `gethostbyaddr()`. This is a `libc` call to resolve an IP address to a hostname. Unfortunately, doing this is quite complicated. When you call `gethostbyaddr()`, a lot of work goes on under the covers. This usually involves making a network call out to the DNS server, and, dangerously, parsing the response.

For clarity (and clarity is a very important part of security), all external APIs used by `vsftpd` are encapsulated within two "system interaction" files, named "`sysutil.c`", and "`sysdeputil.c`" (for the more variable/system dependent calls). This provides a convenient audit point for ascertaining which calls `vsftpd` trusts.

`vsftpd-2.0.0` introduces SSL / TLS support using `OpenSSL`. `OpenSSL` is a massive quantity of code which is essentially parsing complex protocol under the full control of remote malicious clients. SSL / TLS is disabled by default, both at compile time and run time. This forces packagers and administrators to make the decision that they trust the `OpenSSL` library. I personally haven't yet formed an opinion on whether I consider the `OpenSSL` code trustworthy.

#### Summary =====

Be very aware of what APIs and/or programs you are trusting, or you might end up creating a trust relationship which makes your program exploitable -- through no direct fault of your own.