

This document explains the design goals and decisions behind vsftpd.

The importance of a secure design

=====

In a world full of good, careful coders who do not make mistakes, a secure design would not be necessary. After all, in the absence of any programming errors, security would not differ no matter how the program is arranged.

Unfortunately, this is not an ideal world, and coders make plenty of mistakes. Even the careful coders make mistakes. Code auditing is important, and goes some way towards eliminating coding mistakes after the fact. However, we have no guarantee that an audit will catch all the flaws.

So, a secure design acknowledges the possibility of undiscovered flaws, and takes steps to minimise the security impact these flaws can have. An obvious example of something we want to do is to apply the principle of least privilege, which ensure that every part of the program runs with the privilege it needs and no more.

An example of insecure design

=====

Examples of insecure design may be found in most other ftpd's. That's one of the reasons vsftpd has been written. We'll pick on wu-ftp as a specific example, since it is rumoured to run about half of all ftp services.

If I log on to wu-ftp as an anonymous user, a process is run on my behalf to serve my ftp session. Unfortunately, this process typically runs with full root privileges on the remote machine. This means that any security flaw present in parsing the copious ftp protocol will lead to full compromise of that machine. Two concrete examples are the recent wu-ftp format string bug (June 1999), and a buffer overflow dealing with large paths a few months beforehand.

Even OpenBSD's ftpd-BSD had a format string bug leading to remote root compromise of the affected machine, illustrating an earlier point about the requirement for secure design even in the presence of heavy auditing.

Secure design under UNIX

=====

vsftpd is written to run under UNIX-like operating systems, and so its secure

design is constrained by the facilities offered by UNIX. Ideally, UNIX would have a proper security model which would offer fine grained access control to all system interactions (files, network, etc). It doesn't, but it does offer some useful and often overlooked facilities which help us to implement the principle of least privilege:

- Strong inter-process communication facilities

In UNIX, the process is a strongly defined boundary. Different privilege credentials may be assigned to different processes, which are not able to interfere with each other. This is a very basic facility of UNIX.

It makes sense to use this facility to totally separate parts of a program which do not need to be privileged (most) from those parts that do (typically minimal).

The privileged and unprivileged parts of the program then communicate via one of many UNIX IPC mechanisms - perhaps a socketpair or IPC (the former is attractive because UNIX lets you pass file handles over a socket).

The minimal privileged process exercises the "principle of distrust" - it carefully filters what the unprivileged process asks it to do, so that even if the unprivileged process is compromised, it cannot ask the privileged process to do anything we don't want to allow.

- chroot()

chroot() is an often overlooked but useful tool. It can be used very effectively as a damage limitation tool.

Imagine a remotely compromised process which does not run as root, but also does not use chroot(). Now look at what the attacker can do. Amongst the worst items are pilfering of all publicly readable files, and also attempting to execute any publicly executable suid-root programs to try and elevate privilege.

Now imaging the same compromised process with a chroot() to an empty directory. The attackers options to do unpleasant things are substantially diminished.

No, chroot() is not the ideal way to do what we have just accomplished, but it is what we have got to work with. In an ideal environment with fine grained security, we would default to having access to `_no_` files at all, and deliberately not ask for access to any.

- Capabilities (Linux 2.2+)

Like `chroot()`, capabilities are essentially a damage limitation exercise. They are also much less widespread than the other UNIX facilities detailed above. Nonetheless, they warrant mentioning because Linux has them, and they are used in `vsftpd` because that is the primary development platform.

Capabilities split up the all powerful root privilege into lots of sometimes orthogonal privileges. Some of the capabilities represent privileges which are often the basis for requiring a program to run with full root privileges. Examples include `CAP_NET_RAW` (ping, traceroute) and `CAP_NET_BIND_SERVICE` (rlogin).

By using capabilities to ensure we only have the privilege we need (within the somewhat disappointing granularity they offer), we again limit the potential damage of security holes.

Presenting `vsftpd`'s secure design =====

`vsftpd` employs a secure design. The UNIX facilities outlined above are used to good effect. The design decisions taken are as follows:

- 1) All parsing and acting on potentially malicious remote network data is done in a process running as an unprivileged user. Furthermore, this process runs in a `chroot()` jail, ensuring only the ftp files area is accessible.
- 2) Any privileged operations are handled in a privileged parent process. The code for this privileged parent process is as small as possible for safety.
- 3) This same privileged parent process receives requests from the unprivileged child over a socket. All requests are distrusted. Here are example requests:
 - Login request. The child sends username and password. Only if the details are correct does the privileged parent launch a new child with the appropriate user credentials.
 - `chown()` request. The child may request a recently uploaded file gets `chown'ed()` to root for security purposes. The parent is careful to only allow `chown()` to root, and only from files owned by the ftp user.
 - Get privileged socket request. The ftp protocol says we are supposed to emit data connections from port 20. This requires privilege. The privileged parent process creates the privileged socket and passes it to child over the socket.
- 4) This same privileged parent process makes use of capabilities and `chroot()`,

to run with the least privilege required. After login, depending on what options have been selected, the privileged parent dynamically calculates what privileges it requires. In some cases, this amounts to no privilege, and the privileged parent just exits, leaving no part of vsftpd running with privilege.

5) vsftpd-2.0.0 introduces SSL / TLS support using OpenSSL. ALL OpenSSL protocol parsing is performed in a chroot() jail, running under an unprivileged user. This means both pre-authenticated and post-authenticated OpenSSL protocol parsing; it's actually quite hard to do, but vsftpd manages it in the name of being secure. I'm unaware of any other FTP server which supports both SSL / TLS and privilege separation, and gets this right.

Comments on this document are welcomed.